

Isaac Sheets – BAS475 – Assignment4

Isaac Sheets

April 24, 2020

1. (similar) Problem 3.10 in text book Let x_t represent the cardiovascular mortality series "cmort".

```
library(astsa)
data("cmort")
```

- (a) Fit an AR(2) to x_t and report the fitted model and the estimate of the variance of the white noise (σ^2_w).

```
M <- arima(cmort, order=c(2,0,0))
M

##
## Call:
## arima(x = cmort, order = c(2, 0, 0))
##
## Coefficients:
##          ar1      ar2  intercept
##       0.4301  0.4424   88.8538
## s.e.  0.0397  0.0398    1.9407
##
## sigma^2 estimated as 32.37:  log likelihood = -1604.71,  aic = 3217.43
```

Response: $X(t) - 88.8538 = 0.4301(X(t-1) - 88.8538) + 0.4424(X(t-1) - 88.8538) + W_t$; σ^2 estimated as 32.37

- (b) Assuming the fitted model in (a) is the true model, find the forecasts over a four week horizon, x_{n+m} , for $m = 1, 2, 3, 4$, and the corresponding 95% prediction intervals.

```
length(cmort)

## [1] 508

cmort[508] #85.49

## [1] 85.49

cmort[507] #89.43

## [1] 89.43

0.4301*(85.49-88.8538) + 0.4424*(89.43-88.8538) + 88.8538 #1-step ahe
ad prediction

## [1] 87.66194
```

```

0.4301*(87.66194-88.8538) + 0.4424*(85.49-88.8538) + 88.8538      #2-steps ah
ead prediction

## [1] 86.85304

0.4301*(86.85304-88.8538) + 0.4424*(87.66194-88.8538) + 88.8538  #3-steps ah
ead prediction

## [1] 87.46599

0.4301*(87.46599-88.8538) + 0.4424*(86.85304-88.8538) + 88.8538  #4-steps ah
ead prediction

## [1] 87.37177

fore <- predict(M, n.ahead=4) #predict
fore

## $pred
## Time Series:
## Start = c(1979, 41)
## End = c(1979, 44)
## Frequency = 52
## [1] 87.66207 86.85311 87.46615 87.37190
##
## $se
## Time Series:
## Start = c(1979, 41)
## End = c(1979, 44)
## Frequency = 52
## [1] 5.689543 6.193387 7.148343 7.612531

#1-step 95%
c(fore$pred[1]-2*fore$se[1], fore$pred[1]+2*fore$se[1])

## [1] 76.28299 99.04116

#2-steps 95%
c(fore$pred[2]-2*fore$se[2], fore$pred[2]+2*fore$se[2])

## [1] 74.46633 99.23988

#3-steps 95%
c(fore$pred[3]-2*fore$se[3], fore$pred[3]+2*fore$se[3])

## [1] 73.16946 101.76283

#4-steps 95%
c(fore$pred[4]-2*fore$se[4], fore$pred[4]+2*fore$se[4])

## [1] 72.14684 102.59696

```

Response: 1-step ahead prediction: 87.66194; Corresponding confidence interval: 76.28299-99.04116. 2-steps ahead prediction: 86.85304; Corresponding confidence interval: 74.46633-99.23988. 3-steps ahead prediction: 87.46599; Corresponding confidence interval: 73.16946-101.76283. 4-steps ahead prediction: 87.37177; Corresponding confidence interval: 72.14684-102.59696.

2. [(similar) Problem 3.33 in text book] (R) Let's consider fitting an ARIMA(p, d, q) model and generating predictions for the annual global temperature data "gtemp".

```
data("gtemp")
```

- (a) Divide the data set into training and testing data, where the latter data set contains the last 13 data points, i.e., roughly the last 10% of data points, to check the performance of the predictions.

```
length(gtemp)
```

```
## [1] 130
```

```
TESTING <- gtemp[118:130]
```

```
TRAINING <- gtemp[1:117]
```

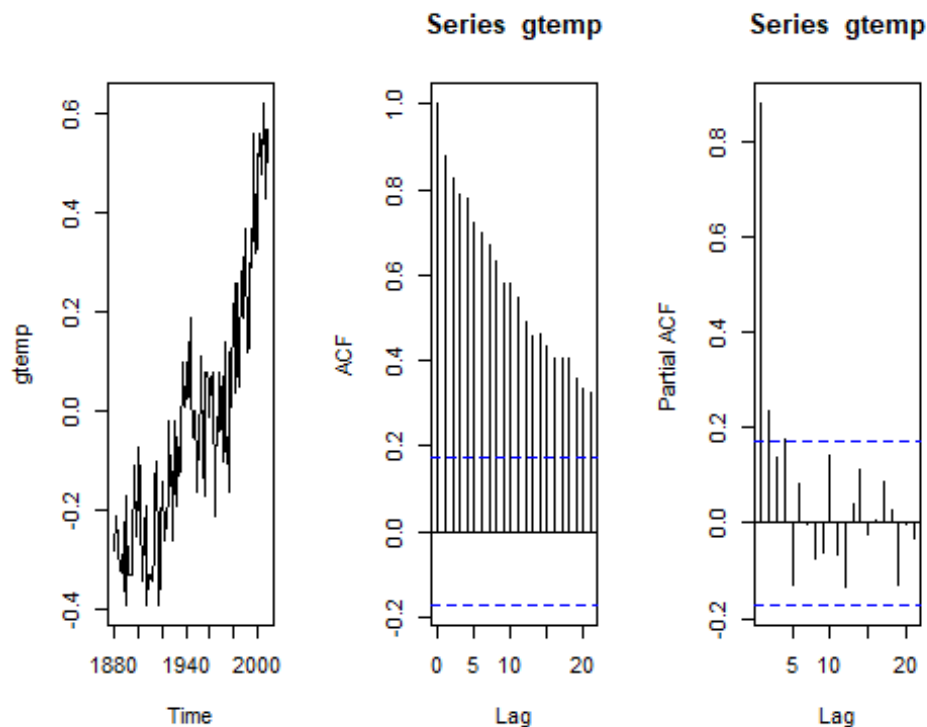
- (b) Check the stationarity of (whole) "gtemp" by looking at the time series plot, ACF and PACF plots.

```
par(mfrow = c(1,3))
```

```
ts.plot(gtemp)
```

```
acf(gtemp)
```


```
pacf(gtemp)
```



Response: The time series is not stationary as the mean changes over time.

(c) Apply the appropriate order of the difference to “gtemp” and make it stationary.

```
gtemp <- diff(gtemp)
plot.ts(gtemp)
```



Response: Since the time series plot was linear, only a single difference needs to be applied here.

(d) Search the optimal ARMA(p, q) model for the differenced training series chosen by AIC/BIC. Try all ARMA(p, q), $p, q = 0, 1, \dots, 5$ and report the selected optimal model(s).

```
n<-length(diff(TRAINING))
P=5
Q=5
crit<-matrix(0,P+1,Q+1)
for (j in 0:P)
{
  for (k in 0:Q)
  {
    dataML<-arima(diff(TRAINING),order=c(j,0,k),method="ML")

    #BIC
    crit[j+1,k+1]<-n*log(dataML$sigma2)+(j+k+1)*log(n)
```

```

}
}

## Warning in arima(diff(TRAINING), order = c(j, 0, k), method = "ML"): possible
## convergence problem: optim gave code = 1

crit

##           [,1]      [,2]      [,3]      [,4]      [,5]      [,6]
## [1,] -512.1347 -533.5280 -535.3289 -531.1026 -529.7051 -525.7690
## [2,] -517.2771 -534.9841 -530.7836 -532.2049 -527.7289 -523.6427
## [3,] -520.5419 -531.3013 -526.8192 -525.2847 -523.1166 -518.9282
## [4,] -530.3512 -527.7152 -528.1880 -523.6153 -519.1286 -519.3701
## [5,] -525.7062 -530.1493 -527.6085 -524.0235 -519.2679 -514.4944
## [6,] -526.1900 -522.0180 -523.6455 -519.2660 -514.2593 -514.2610

min(crit)

## [1] -535.3289

n<-length(diff(TRAINING))
P=5
Q=5
crit<-matrix(0,P+1,Q+1)
for (j in 0:P)
{
  for (k in 0:Q)
  {
    dataML<-arima(diff(TRAINING),order=c(j,0,k),method="ML")

    #AIC
    crit[j+1,k+1]<-n*log(dataML$sigma2)+2*(j+k+1)

  }
}

## Warning in arima(diff(TRAINING), order = c(j, 0, k), method = "ML"): possible
## convergence problem: optim gave code = 1

crit

##           [,1]      [,2]      [,3]      [,4]      [,5]      [,6]
## [1,] -514.8883 -539.0351 -543.5896 -542.1170 -543.4731 -542.2905
## [2,] -522.7843 -543.2449 -541.7979 -545.9729 -544.2505 -542.9178
## [3,] -528.8027 -542.3156 -540.5871 -541.8062 -542.3917 -540.9569
## [4,] -541.3655 -541.4831 -544.7095 -542.8904 -541.1573 -544.1524
## [5,] -539.4742 -546.6708 -546.8836 -546.0522 -544.0502 -542.0303
## [6,] -542.7115 -541.2932 -545.6742 -544.0483 -541.7952 -544.5505

min(crit)

```

```
## [1] -546.8836
```

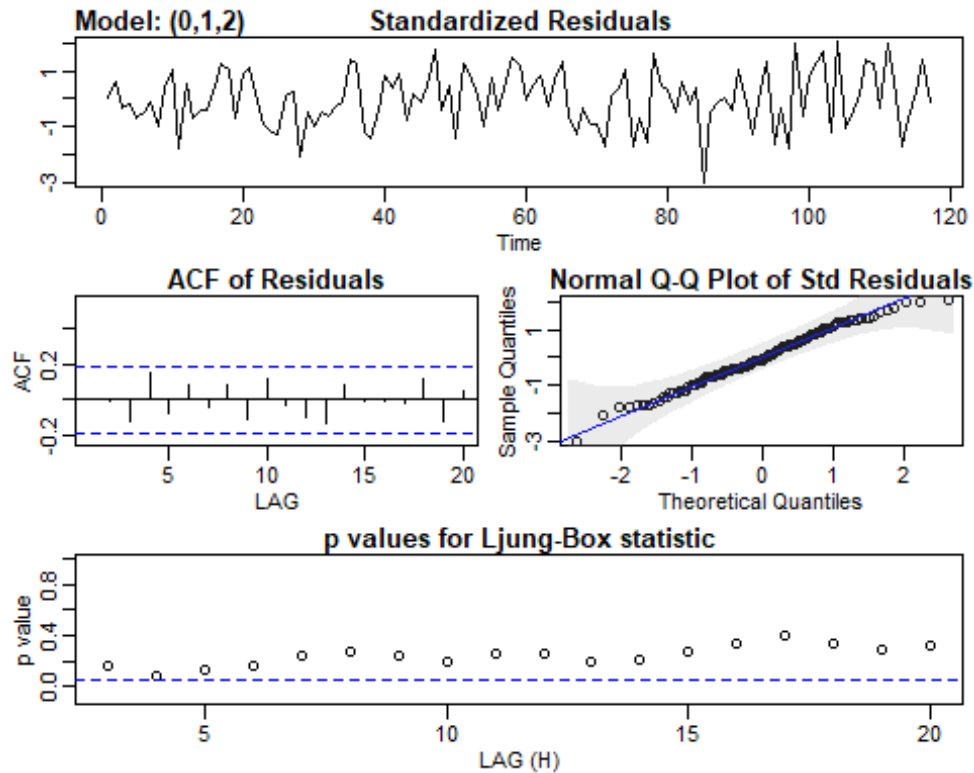
Response: BIC suggest an MA(2) model; AIC also suggests an ARMA(4,2) model.

(e) Fit the optimal ARIMA(p, d, q) model to (no differenced) training data, where p, q are selected from (c) and d is selected from (b).

(f) Check the diagnostics plots and comment on each plot.

```
M <- sarima(TRAINING, 0,1,2)
```

```
## initial  value -2.227967
## iter    2 value -2.325083
## iter    3 value -2.365718
## iter    4 value -2.367238
## iter    5 value -2.367902
## iter    6 value -2.368766
## iter    7 value -2.368774
## iter    8 value -2.368775
## iter    8 value -2.368775
## final   value -2.368775
## converged
## initial  value -2.365734
## iter    2 value -2.365746
## iter    3 value -2.365767
## iter    4 value -2.365775
## iter    4 value -2.365775
## iter    4 value -2.365775
## final   value -2.365775
## converged
```



M

```
## $fit
##
## Call:
## stats::arima(x = xdata, order = c(p, d, q), seasonal = list(order = c(P, D
,
##     Q), period = S), xreg = constant, transform.pars = trans, fixed = fixe
d,
##     optim.control = list(trace = trc, REPORT = 1, reltol = tol))
##
## Coefficients:
##          ma1      ma2  constant
##      -0.5323 -0.2212   0.0048
## s.e.   0.0848  0.0822   0.0022
##
## sigma^2 estimated as 0.008758:  log likelihood = 109.83,  aic = -211.67
##
## $degrees_of_freedom
## [1] 113
##
## $ttable
##      Estimate      SE t.value p.value
## ma1      -0.5323 0.0848 -6.2795  0.0000
## ma2      -0.2212 0.0822 -2.6913  0.0082
## constant   0.0048 0.0022  2.1718  0.0320
##
```

```
## $AIC
## [1] -1.824708
##
## $AICc
## [1] -1.82286
##
## $BIC
## [1] -1.729756

N <- sarima(TRAINING, 4,1,2)

## initial  value -2.214331
## iter    2 value -2.350495
## iter    3 value -2.369009
## iter    4 value -2.372395
## iter    5 value -2.376101
## iter    6 value -2.378948
## iter    7 value -2.381024
## iter    8 value -2.383901
## iter    9 value -2.387263
## iter   10 value -2.392602
## iter   11 value -2.399314
## iter   12 value -2.401882
## iter   13 value -2.404254
## iter   14 value -2.407097
## iter   15 value -2.410847
## iter   16 value -2.415446
## iter   17 value -2.415682
## iter   18 value -2.427094
## iter   19 value -2.430705
## iter   20 value -2.433987
## iter   21 value -2.434603
## iter   22 value -2.435243
## iter   22 value -2.435243
## iter   23 value -2.439460
## iter   24 value -2.439819
## iter   24 value -2.439819
## iter   25 value -2.439933
## iter   26 value -2.440027
## iter   27 value -2.440224
## iter   28 value -2.440307
## iter   29 value -2.440492
## iter   30 value -2.440573
## iter   31 value -2.440736
## iter   32 value -2.440817
## iter   33 value -2.440956
## iter   34 value -2.441039
## iter   35 value -2.441155
## iter   36 value -2.441240
## iter   37 value -2.441334
```




```
## iter 38 value -2.441420
## iter 39 value -2.441494
## iter 40 value -2.441581
## iter 40 value -2.441581
## iter 41 value -2.441638
## iter 42 value -2.441723
## iter 42 value -2.441723
## iter 43 value -2.441765
## iter 44 value -2.441849
## iter 44 value -2.441849
## iter 45 value -2.441881
## iter 46 value -2.441960
## iter 47 value -2.441960
## iter 48 value -2.441984
## iter 49 value -2.442061
## iter 50 value -2.442064
## iter 51 value -2.442105
## iter 52 value -2.442264
## iter 52 value -2.442264
## iter 53 value -2.442467
## iter 54 value -2.442472
## iter 54 value -2.442472
## iter 55 value -2.442502
## iter 56 value -2.442605
## iter 56 value -2.442605
## iter 57 value -2.442647
## iter 58 value -2.442666
## iter 59 value -2.442668
## iter 60 value -2.442677
## iter 61 value -2.442690
## iter 61 value -2.442690
## iter 62 value -2.442696
## iter 63 value -2.442704
## iter 63 value -2.442704
## iter 64 value -2.442712
## iter 65 value -2.442722
## iter 65 value -2.442722
## iter 66 value -2.442727
## iter 67 value -2.442736
## iter 67 value -2.442736
## iter 68 value -2.442742
## iter 69 value -2.442750
## iter 69 value -2.442750
## iter 70 value -2.442756
## iter 71 value -2.442764
## iter 71 value -2.442764
## iter 72 value -2.442770
## iter 73 value -2.442777
## iter 73 value -2.442777
## iter 74 value -2.442784
```

```
## iter 75 value -2.442791
## iter 75 value -2.442791
## iter 76 value -2.442797
## iter 77 value -2.442804
## iter 77 value -2.442804
## iter 78 value -2.442810
## iter 79 value -2.442816
## iter 79 value -2.442816
## iter 80 value -2.442823
## iter 81 value -2.442829
## iter 81 value -2.442829
## iter 82 value -2.442835
## iter 83 value -2.442841
## iter 83 value -2.442841
## iter 84 value -2.442848
## iter 85 value -2.442853
## iter 85 value -2.442853
## iter 86 value -2.442860
## iter 87 value -2.442865
## iter 87 value -2.442865
## iter 88 value -2.442872
## iter 89 value -2.442877
## iter 89 value -2.442877
## iter 90 value -2.442884
## iter 91 value -2.442889
## iter 91 value -2.442889
## iter 92 value -2.442896
## iter 93 value -2.442901
## iter 93 value -2.442901
## iter 94 value -2.442908
## iter 95 value -2.442913
## iter 95 value -2.442913
## iter 96 value -2.442919
## iter 97 value -2.442924
## iter 97 value -2.442924
## iter 98 value -2.442931
## iter 99 value -2.442936
## iter 99 value -2.442936
## iter 100 value -2.442942
## final value -2.442942
## stopped after 100 iterations
## initial value -2.227967
## iter 2 value -2.347959
## iter 3 value -2.377795
## iter 4 value -2.379996
## iter 5 value -2.383626
## iter 6 value -2.386832
## iter 7 value -2.388224
## iter 8 value -2.391012
## iter 9 value -2.392302
```

```

## iter 10 value -2.393382
## iter 11 value -2.395564
## iter 12 value -2.395738
## iter 13 value -2.397841
## iter 14 value -2.400238
## iter 15 value -2.401668
## iter 16 value -2.402203
## iter 17 value -2.402228
## iter 18 value -2.402267
## iter 19 value -2.402275
## iter 20 value -2.402278
## iter 21 value -2.402295
## iter 22 value -2.402325
## iter 23 value -2.402351
## iter 24 value -2.402358
## iter 25 value -2.402359
## iter 26 value -2.402359
## final value -2.402359
## converged

```



N

```

## $fit
##
## Call:
## stats::arima(x = xdata, order = c(p, d, q), seasonal = list(order = c(P, D

```

```
,
##      Q), period = S), xreg = constant, transform.pars = trans, fixed = fixe
d,
##      optim.control = list(trace = trc, REPORT = 1, reltol = tol))
##
## Coefficients:
##           ar1      ar2      ar3      ar4      ma1      ma2  constant
##          -0.1902  0.3022 -0.0115  0.2499 -0.3026 -0.6974   0.0046
## s.e.      0.2470  0.1435  0.0972  0.1000  0.2537  0.2530   0.0006
##
## sigma^2 estimated as 0.007945:  log likelihood = 114.08,  aic = -212.15
##
## $degrees_of_freedom
## [1] 109
##
## $ttable
##      Estimate      SE t.value p.value
## ar1      -0.1902 0.2470 -0.7700  0.4430
## ar2       0.3022 0.1435  2.1066  0.0374
## ar3      -0.0115 0.0972 -0.1178  0.9065
## ar4       0.2499 0.1000  2.4998  0.0139
## ma1      -0.3026 0.2537 -1.1925  0.2357
## ma2      -0.6974 0.2530 -2.7562  0.0069
## constant  0.0046 0.0006  7.6481  0.0000
##
## $AIC
## [1] -1.82891
##
## $AICc
## [1] -1.81997
##
## $BIC
## [1] -1.639007
```

Response:

ARIMA(0,1,2):

$X(t) - 0.0048 = W(t) - 0.5323W(t-1) - 0.2212W(t-2)$; σ^2 estimated as 0.008758

ARIMA(4,1,2):

$X(t) - 0.0046 = -0.1902(X(t-1) - 0.0046) + 0.3022(X(t-2) - 0.0046) - 0.0115(X(t-3) - 0.0046) + 0.2499(X(t-4) - 0.0046) + W(t) - 0.3026W(t-1) - 0.6974W(t-2)$; σ^2 estimated as 0.007945

The diagnostics for the ARIMA(0,1,2) look great. All points on the Standardized Residuals plot appear to be within the ± 3 interval. No ACF's of Residuals appear to be significant. The QQ plot looks great, and all p-values for the Ljung-Box test are above 0.05.


The diagnostics for the ARIMA(4,1,2) also look great. All points on the Standardized Residuals plot appear to be within the ± 3 interval. No ACF's of Residuals appear to be significant. The QQ plot looks great, and all p-values for the Ljung-Box test are above 0.05.

- (g) After the optimal model(s) is fitted, forecast (with respect to "gtemp" not the differenced series) the next 13 years which belongs to the testing set. Plot the predictions, 95% prediction intervals with the observed data in one graph. Assess the performance of the predictions, i.e., observe if the actual testing data points are included in the prediction intervals.

```
par(mfrow = c(1,1))
sarima.for(TRAINING, n.ahead = 13, 0,1,2)

## $pred
## Time Series:
## Start = 118
## End = 130
## Frequency = 1
## [1] 0.2693195 0.2758779 0.2807009 0.2855238 0.2903468 0.2951697 0.2999927
## [8] 0.3048156 0.3096386 0.3144615 0.3192845 0.3241074 0.3289303
##
## $se
## Time Series:
## Start = 118
## End = 130
## Frequency = 1
## [1] 0.09358163 0.10331195 0.10585581 0.10833995 0.11076840 0.11314474
## [7] 0.11547218 0.11775363 0.11999171 0.12218881 0.12434709 0.12646854
## [13] 0.12855498


points(118:130,TESTING, col="blue", pch=20)
lines(118:130, TESTING)
```



```
par(mfrow = c(1,1))
sarima.for(TRAINING, n.ahead = 13, 4,1,2)

## $pred
## Time Series:
## Start = 118
## End = 130
## Frequency = 1
## [1] 0.2404997 0.2332837 0.2586037 0.2351919 0.2380202 0.2313221 0.2430551
## [8] 0.2359254 0.2446202 0.2420125 0.2511590 0.2497590 0.2580014
##
## $se
## Time Series:
## Start = 118
## End = 130
## Frequency = 1
## [1] 0.08949941 0.10068152 0.10254911 0.10307143 0.10653665 0.10704765
## [7] 0.10772170 0.10779898 0.10831371 0.10834159 0.10854430 0.10854417
## [13] 0.10867296

points(118:130,TESTING, col="blue", pch=20)
lines(118:130, TESTING)
```




Response: There are 4 points outside of the 95% confidence interval in the ARIMA(0,1,2) model, and there are 9 points outside of the 95% confidence interval for the ARIMA(4,1,2). It is clear that ARIMA(0,1,2) is the better performing model.

3. [Problem 3.36 in text book] (R) One of the remarkable technological developments in the computer industry has been the ability to store information densely on a hard drive. In addition, the cost of storage has steadily declined causing problems of too much data as opposed to big data. The data set for this assignment is “cpg”, which consists of the median annual retail price per GB of hard drives, say *ct*, taken from a sample of manufacturers from 1980 to 2008.

```
data(cpg)
```

- (a) Plot *ct* and describe what you see.


```
plot.ts(cpg)
```



Response: There appears to be an exponential (quadratic) decrease.

- (b) It seems that the curve c_t versus t behaves like $c_t \approx \alpha e^{\beta \cdot t}$. Then let's fit a linear regression of $\log c_t$ on t and plot the fitted line to compare it to the logged data. Comment on what you observe. [hint: When you fit a linear regression, use `lm(log(cpg)~time(cpg))`. After fitting a regression model, first plot `thelog(cpg)` and use `abline()` function to add the fitted line]

```
K <- lm(log(cpg)~time(cpg))
plot.ts(log(cpg))
abline(K, col="blue")
```


Response: It appears that it changed the exponential decrease into a linear decrease.

(c) Inspect the residuals of the linear regression fit by generating the time series plot, ACF and PACF plot. Visually decide which ARMA model is appropriate for the residual.

`K$residuals`


```
##          1          2          3          4          5          6
## -1.77156094 -0.86080012 -0.40201622 -0.21283425  0.28263122  0.05521491
##          7          8          9         10         11         12
##  0.47195722  0.36388767  1.13128685  0.33007012  0.66383332  0.68929516
##          13         14         15         16         17         18
##  0.58122560  0.42186275  0.91320790  0.29238409 -0.04463736  0.04725744
##          19         20         21         22         23         24
## -0.31053798 -0.39840482 -0.48119657 -0.76816142 -0.71782497 -0.26173946
##          25         26         27         28         29
## -0.39922290 -0.04854598  0.30390936  0.44715423 -0.31769486
```

`plot.ts(K$residuals)`

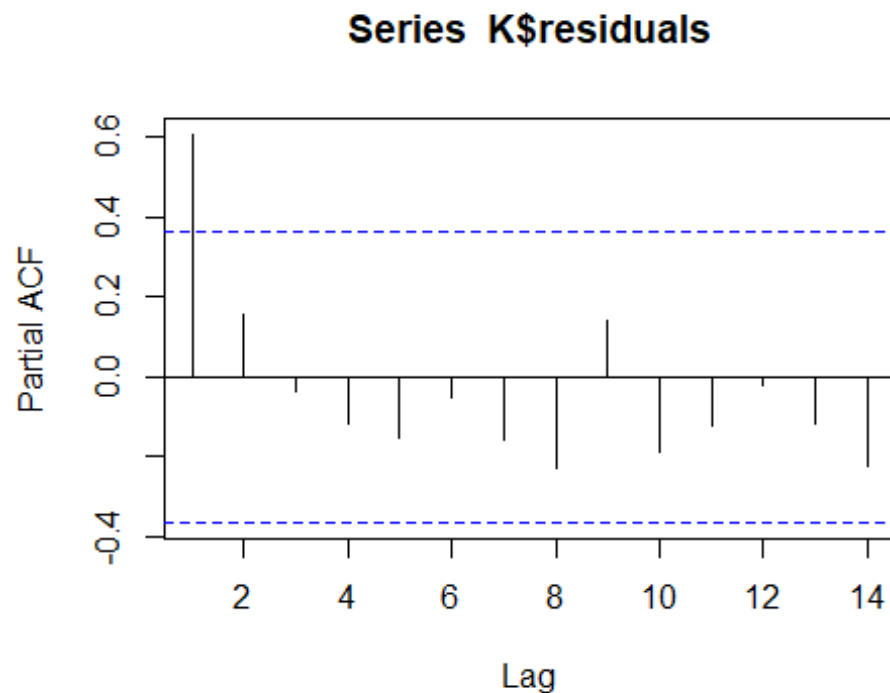


```
acf(K$residuals)
```

Series K\$residuals



```
pacf(K$residuals)
```



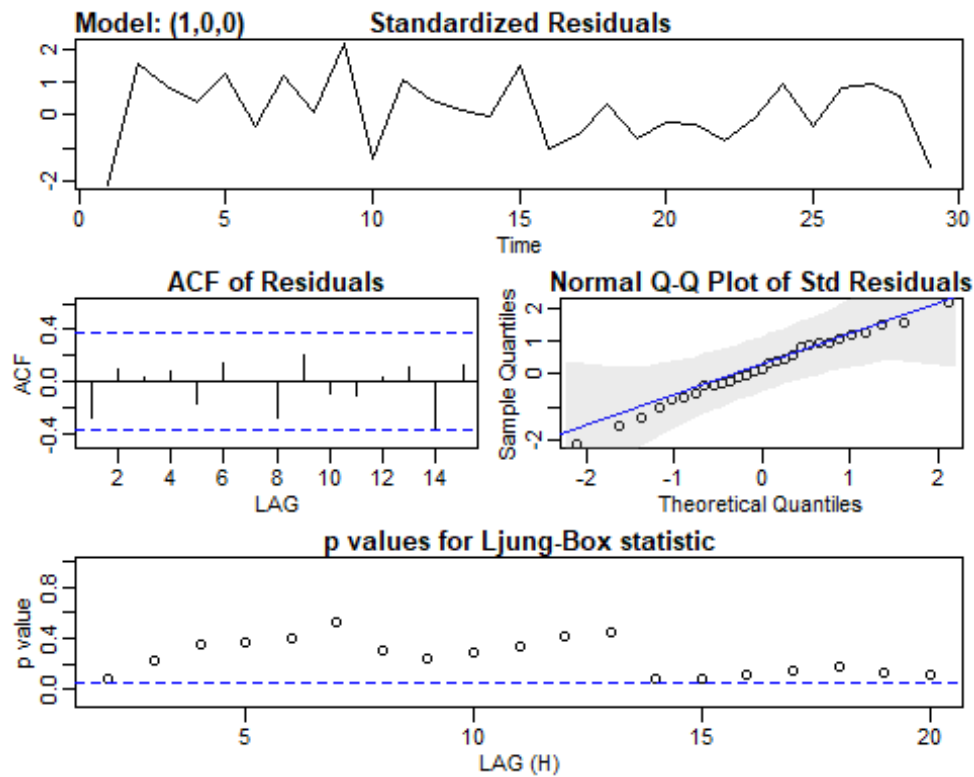
Response: ACF tails off and the PACF cuts off at lag=1. Therefore, we should use the AR(1) model.

(d) Fit the model chosen from (c) to the residuals and report the fitted model.

```
G <- sarima(K$residuals, 1,0,0)
```

```
## initial value -0.669056
## iter 2 value -0.968880
## iter 3 value -0.991232
## iter 4 value -1.051693
## iter 5 value -1.053940
## iter 6 value -1.061132
## iter 7 value -1.061236
## iter 8 value -1.061270
## iter 9 value -1.061293
## iter 10 value -1.061293
## iter 10 value -1.061293
## iter 10 value -1.061293
## final value -1.061293
## converged
## initial value -0.809581
## iter 2 value -0.868173
## iter 3 value -0.873691
## iter 4 value -0.876061
## iter 5 value -0.876326
## iter 6 value -0.876361
## iter 7 value -0.876367
```

```
## iter    8 value -0.876370
## iter    9 value -0.876370
## iter   10 value -0.876370
## iter   10 value -0.876370
## iter   10 value -0.876370
## final  value -0.876370
## converged
```



G

```
## $fit
##
## Call:
## stats::arima(x = xdata, order = c(p, d, q), seasonal = list(order = c(P, D
,
##     Q), period = S), xreg = xmean, include.mean = FALSE, transform.pars =
trans,
##     fixed = fixed, optim.control = list(trace = trc, REPORT = 1, reltol =
tol))
##
## Coefficients:
##          ar1      xmean
##      0.8239  -0.2548
## s.e.  0.1253   0.4096
##
## sigma^2 estimated as 0.1666:  log likelihood = -15.73,  aic = 37.47
##
```

```
## $degrees_of_freedom
## [1] 27
##
## $ttable
##      Estimate      SE t.value p.value
## ar1      0.8239 0.1253  6.5763  0.0000
## xmean   -0.2548 0.4096 -0.6221  0.5391
##
## $AIC
## [1] 1.292033
##
## $AICc
## [1] 1.307948
##
## $BIC
## [1] 1.433478

summary(K)

##
## Call:
## lm(formula = log(cpg) ~ time(cpg))
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.77156 -0.39840  0.04726  0.42186  1.13129
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 1172.49431    27.57793   42.52  <2e-16 ***
## time(cpg)    -0.58508     0.01383  -42.30  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.6231 on 27 degrees of freedom
## Multiple R-squared:  0.9851, Adjusted R-squared:  0.9846
## F-statistic: 1790 on 1 and 27 DF,  p-value: < 2.2e-16
```

Response: $X(t) = 0.8239X(t-1) + W_t$; σ^2 estimated as 0.1666

(e) Report the final model with respect to $\log(ct)$. Combine the results of a linear regression part and time series part.

Response:

**$Y(t) - 1172.49431 + 0.58508t = 0.82(Y(t-1) - 1172.49431 + 0.58508(t-1)) + W(t)$;
 σ^2 estimated as 0.1666**