

BAS 474: Final Coding Exam

Isaac Sheets

Due: April 26, 2020

- Do not discuss this work with anyone. If you have any questions, you should consult your instructor.
- Your instructors will not answer questions on Sunday.
- For questions without tuning parameters given, you don't have to include `tuneGrid` in the `train` function. In this way some default parameters will be used. You can also set the parameters up by following examples from slides or assignments.

Problem 1. Predictive Model For Tip Percentage

The TIPS dataset in the `regclass` package contains information about tips a waiter received over a period of a few months working in one restaurant. The waiter collected information in 8 variables on 244 waiting services (observations). Your task is to build a predictive model for `TipPercentage`. The variables available in the data are:

- `TipPercentage`: a numeric vector, the tip written as a percentage (0-100) of the total bill
- `Bill`: a numeric vector, the bill amount (dollars)
- `Tip`: a numeric vector, the tip amount (dollars)
- `Gender`: a factor with levels Female Male, gender of the payer of the bill
- `Smoker`: a factor with levels No Yes, whether the party included smokers
- `Weekday`: a factor with levels Friday Saturday Sunday Thursday, day of the week
- `Time`: a factor with levels Day Night, rough time of day
- `PartySize`: a numeric vector, number of people in party

Do not use the `Tip` variable as predictor (you can remove this variable by nulling it out).

Split the data into 80% training rows with the remaining rows being the holdout (use `set.seed(474)` on the same line as the required `sample` command).

Use `set.seed(474)` everywhere randomness is infused in the training process.

Use 5-fold cross-validation to estimate the generalization error.

```
data(TIPS)
TIPS$Tip <- NULL
set.seed(474); train.rows <- sample(nrow(TIPS), 0.8*nrow(TIPS))
TRAIN <- TIPS[train.rows,];
```

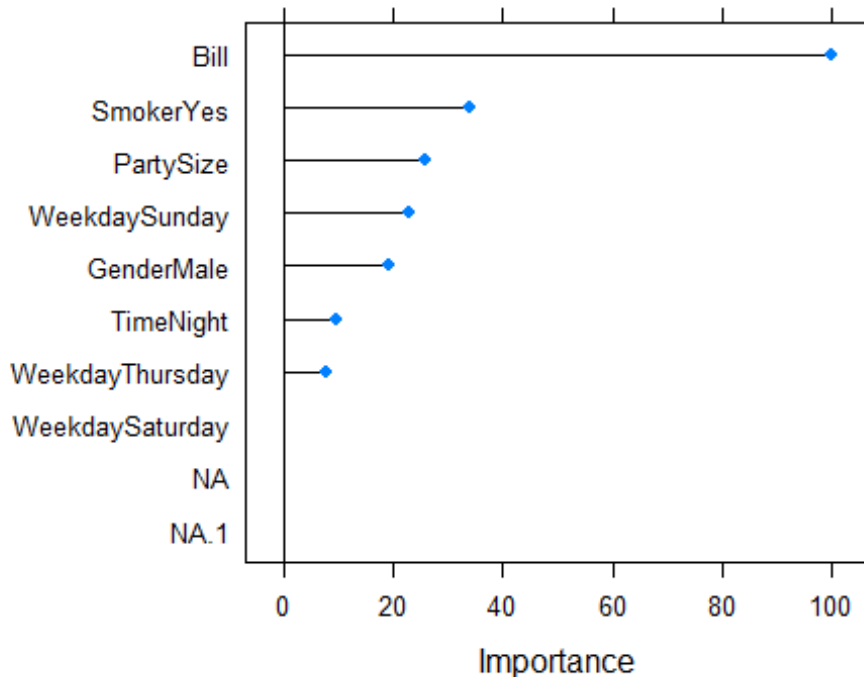
```
HOLDOUT <- TIPS[-train.rows,]
fitControl <- trainControl(method="cv", number=5)
```

- (a) Build a vanilla linear regression model using the training data to predict TipPercentage. Report the estimated generalization RMSE and the RMSE on the holdout sample. Report the variable importance plot, and comment on which predictors appear most important for predicting TipPercentage.

```
set.seed(474); GLM <- train(TipPercentage~., data=TRAIN, method="glm",
                             trControl=fitControl, preProc=c("center", "scale"
))
GLM$results
## parameter      RMSE Rsquared      MAE  RMSESD RsquaredSD      MAESD
## 1      none 5.904779 0.1390757 4.134148 2.116873 0.07076399 0.6548951

set.seed(474); classification.glm <- predict(GLM, newdata=HOLDOUT)
RMSE(HOLDOUT$TipPercentage, classification.glm)
## [1] 4.620693

IMP <- varImp(GLM)$importance
IMP$Variable <- rownames(IMP)
IMP <- IMP[order(IMP$Overall, decreasing=TRUE),]
head(IMP)
##           Overall      Variable
## Bill      100.000000      Bill
## SmokerYes   34.081084   SmokerYes
## PartySize   25.785500   PartySize
## WeekdaySunday 23.071185 WeekdaySunday
## GenderMale  19.193405   GenderMale
## TimeNight   9.664646   TimeNight
plot(varImp(GLM), top=10)
```



RESPONSE: The estimated generalization RMSE is 5.527699. The RMSE of the holdout sample is 5.63277. The predictors that appear to be most important for predicting TipPercentage are Bill, PartySize, and SmokersYes.

- (b) Build a regularized multiple linear regression model using the training data to predict TipPercentage. Audition alpha values along the sequence 0, 0.1, 0.2, ..., 0.9, 1 and lambda of 10 raised to the -4, -3.5, ..., 1.5, 2 powers. Report the estimated generalization RMSE and the RMSE on the holdout sample. Report the variable importance plot, and comment on which predictors appear most important for predicting TipPercentage.

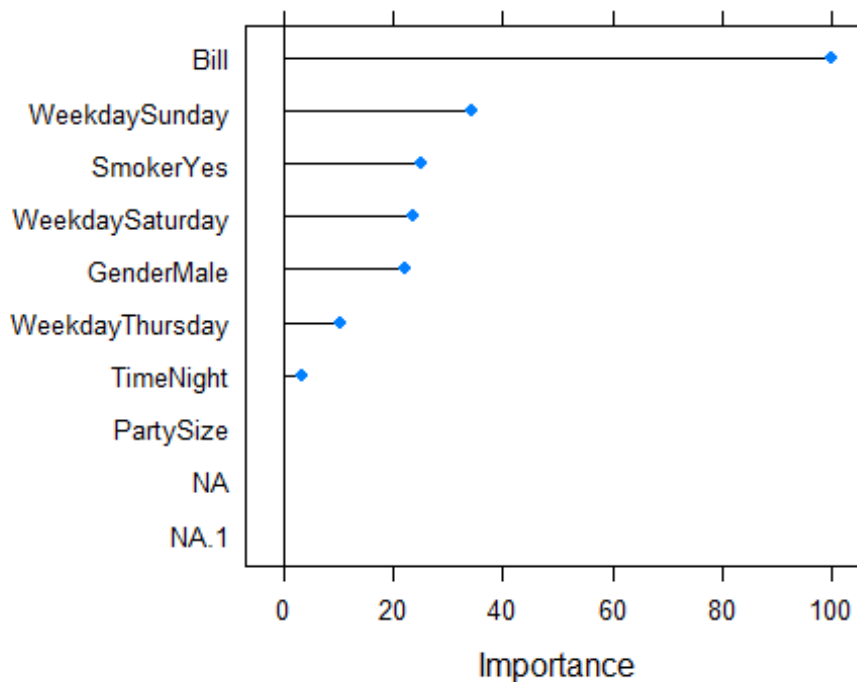
```
glmnetGrid <- expand.grid(alpha = seq(0,1, by=0.1), lambda = 10^seq(-4,2, by=
0.5))
set.seed(474); GLMnet <- train(TipPercentage~., data=TRAIN, method="glmnet",
tuneGrid=glmnetGrid,
                                trControl=fitControl, preProc=c("center", "scale
"))
## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info = trainI
nfo, :
## There were missing values in resampled performance measures.
GLMnet$results[rownames(GLMnet$bestTune),]
##   alpha  lambda  RMSE Rsquared  MAE  RMSESD RsquaredSD  MAESD
## 10      0 3.162278 5.813812 0.1323445 3.98199 2.257199 0.06256246 0.6354841

set.seed(474); classification.glmnet <- predict(GLMnet,newdata=HOLDOUT)
RMSE(HOLDOUT$TipPercentage, classification.glmnet)
## [1] 4.579815
```

```

IMP <- varImp(GLMnet)$importance
IMP$Variable <- rownames(IMP)
IMP <- IMP[order(IMP$Overall, decreasing=TRUE),]
head(IMP)
##           Overall      Variable
## Bill          100.00000      Bill
## WeekdaySunday  34.20049 WeekdaySunday
## SmokerYes      25.07473   SmokerYes
## WeekdaySaturday 23.80801 WeekdaySaturday
## GenderMale     22.26291   GenderMale
## WeekdayThursday 10.49329 WeekdayThursday
plot(varImp(GLMnet), top=10)

```



Response: The estimated generalization RMSE is 5.46346. The RMSE of the holdout sample is 5.607496. The predictors that appear to be most important for predicting TipPercentage are Bill, PartySize, SmokersYes, WeekdaySunday, and WeekdaySaturday.

- (c) Build a KNN model using your training data to predict TipPercentage. Audition k values of 1-30. Report the estimated generalization RMSE and the RMSE on the holdout sample. Report the variable importance plot, and comment on which predictors appear most important for predicting TipPercentage.

```

knnGrid <- expand.grid(k=1:30)
set.seed(474); KNN <- train(TipPercentage~., data=TRAIN, method='knn', preProc
= c("center", "scale"),

```


```

trControl=fitControl, tuneGrid=knnGrid)
KNN$results[which(KNN$results$k == as.numeric(KNN$bestTune) ),]
##      k      RMSE Rsquared      MAE      RMSESD RsquaredSD      MAESD
## 14 14 5.742774 0.1371878 4.02328 2.165275 0.07539694 0.6360195

set.seed(474); classification.knn <- predict(KNN,newdata=HOLDOUT)
RMSE(HOLDOUT$TipPercentage, classification.knn)
## [1] 4.489457

IMP <- varImp(KNN)$importance
IMP$Variable <- rownames(IMP)
IMP <- IMP[order(IMP$Overall,decreasing=TRUE),]
head(IMP)
##              Overall Variable
## Bill      100.0000000      Bill
## PartySize   7.9935800 PartySize
## Gender       4.1750099      Gender
## Smoker       0.5760554      Smoker
## Weekday      0.4241152      Weekday
## Time         0.0000000        Time
plot(varImp(KNN), top=10)

```



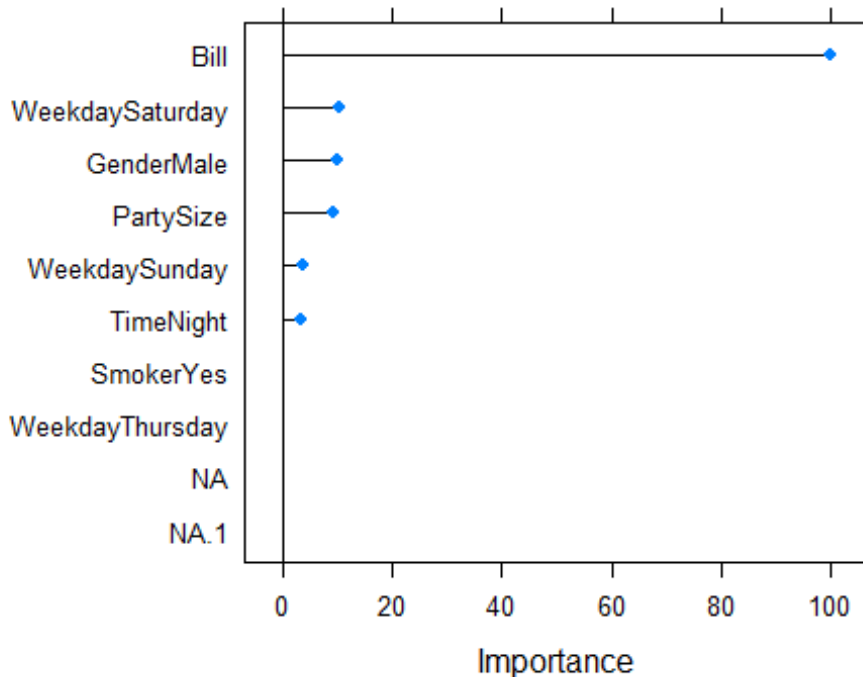
RESPONSE: The estimated generalization RMSE is 5.566491. The RMSE on the holdout sample is 5.513077. It appears that the only important predictor for predicting TipPercentage is Bill.

- (d) Build a regression tree model using your training data to predict TipPercentage. Report the estimated generalization RMSE and the RMSE on the holdout sample. Report the variable importance plot, and comment on which predictors appear most important for predicting TipPercentage.

```
set.seed(474); TREE <- train(TipPercentage~., data=TRAIN, method="rpart", trControl=fitControl,
                             preProc=c("center", "scale"))
## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info = trainInfo, :
## There were missing values in resampled performance measures.
TREE$results[rownames(TREE$bestTune),]
##           cp      RMSE  Rsquared      MAE   RMSESD RsquaredSD      MAESD
## 1 0.03589843 6.132361 0.1241258 4.132534 1.508779 0.1275679 0.3703102

set.seed(474); classification.tree <- predict(TREE, newdata=HOLDOUT)
RMSE(HOLDOUT$TipPercentage, classification.tree)
## [1] 4.935643

IMP <- varImp(TREE)$importance
IMP$Variable <- rownames(IMP)
IMP <- IMP[order(IMP$Overall, decreasing=TRUE),]
head(IMP)
##           Overall      Variable
## Bill          100.000000      Bill
## WeekdaySaturday 10.529947 WeekdaySaturday
## GenderMale      9.948801      GenderMale
## PartySize       9.429961      PartySize
## WeekdaySunday   3.977997      WeekdaySunday
## TimeNight       3.441621      TimeNight
plot(varImp(TREE), top=10)
```



RESPONSE: The estimated generalization RMSE is 5.76658. The RMSE on the holdout sample is 5.75445. It appears that the only important predictor for predicting TipPercentage is Bill.

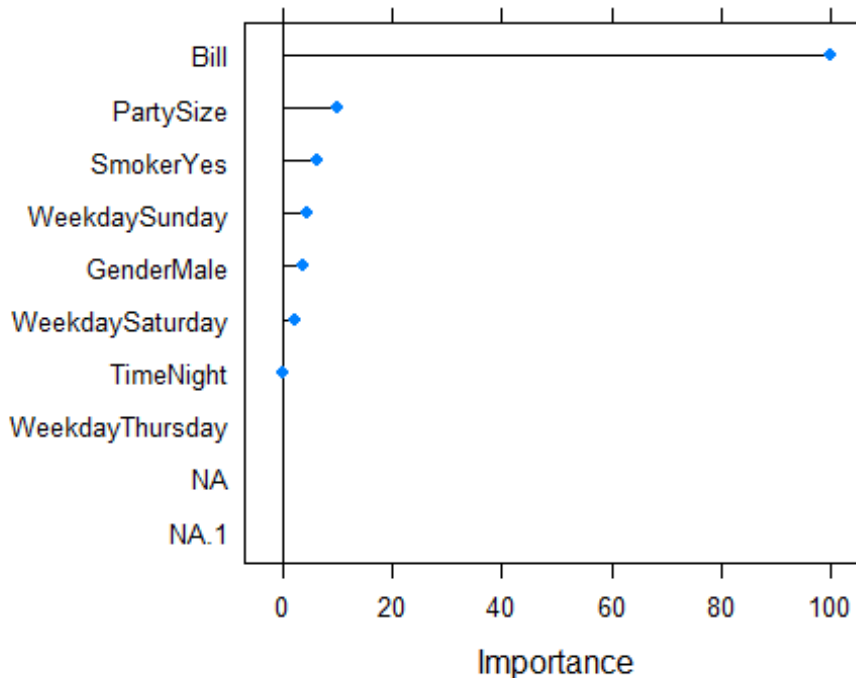
- (e) Build a random forest model using your training data to predict TipPercentage. Audition values of `mtry` of 1-5. Report the estimated generalization RMSE and the RMSE on the holdout sample. Report the variable importance plot, and comment on which predictors appear most important for predicting TipPercentage.

```
forestGrid <- expand.grid(mtry=seq(5))
set.seed(474); FOREST <- train(TipPercentage~.,data=TRAIN,method="rf",tuneGrid=forestGrid,
                               trControl=fitControl,preProc=c("center","scale"))
FOREST$results[rownames(FOREST$bestTune),]
##   mtry    RMSE Rsquared    MAE  RMSESD RsquaredSD    MAESD
## 2     2 5.832596 0.1560506 3.979021 1.918469 0.1393561 0.5258356

set.seed(474); classification.forest <- predict(FOREST,newdata=HOLDOUT)
RMSE(HOLDOUT$TipPercentage, classification.forest)
## [1] 4.881135

IMP <- varImp(FOREST)$importance
IMP$Variable <- rownames(IMP)
IMP <- IMP[order(IMP$Overall,decreasing=TRUE),]
head(IMP)
##           Overall      Variable
```

```
## Bill      100.000000      Bill
## PartySize  10.003010      PartySize
## SmokerYes   6.401454      SmokerYes
## WeekdaySunday  4.477199      WeekdaySunday
## GenderMale   3.757691      GenderMale
## WeekdaySaturday  2.200660      WeekdaySaturday
plot(varImp(FOREST), top=10)
```



RESPONSE: The estimated generalization RMSE is 5.678023. The RMSE on the holdout sample is 5.894915. It appears that the only important predictor for predicting TipPercentage is Bill.

- (f) Build a gradient boosted tree model using the training data to predict TipPercentage. Report the estimated generalization RMSE and the RMSE on the holdout sample. Report the variable importance plot, and comment on which predictors appear most important for predicting TipPercentage.

```
set.seed(474); GBM <- train(TipPercentage~., data=TRAIN, method="gbm",
                           trControl=fitControl, preProc=c("center", "scale"),
                           verbose=FALSE)
GBM$results[rownames(GBM$bestTune),]
## shrinkage interaction.depth n.minobsinnode n.trees      RMSE  Rsquared
## 7      0.1              3          10      50 5.709391 0.1757444
##      MAE  RMSESD RsquaredSD      MAESD
## 7 4.050114 1.99549  0.1020458 0.5408795

set.seed(474); classification.gbm <- predict(GBM, newdata=HOLDOUT)
```

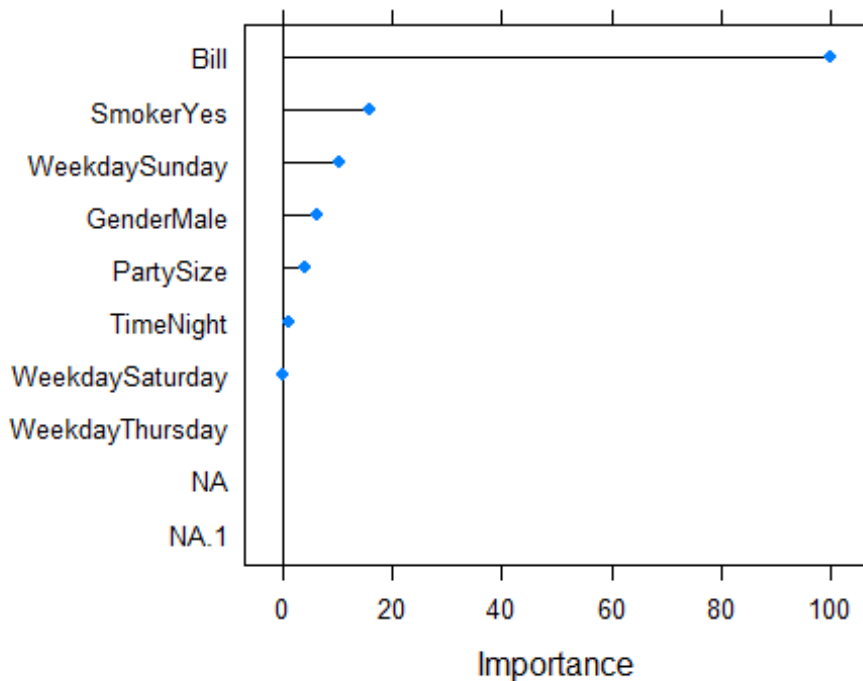


```

RMSE(HOLDOUT$TipPercentage, classification.gbm)
## [1] 4.874882

IMP <- varImp(gbm)$importance
IMP$Variable <- rownames(IMP)
IMP <- IMP[order(IMP$Overall, decreasing=TRUE),]
head(IMP)
##           Overall      Variable
## Bill          100.000000      Bill
## SmokerYes      15.790421    SmokerYes
## WeekdaySunday  10.570905 WeekdaySunday
## GenderMale      6.291911    GenderMale
## PartySize       4.035395    PartySize
## TimeNight       1.093058    TimeNight
plot(varImp(gbm), top=10)

```



RESPONSE: The estimated generalization RMSE is 5.48649. The RMSE on the holdout sample is 5.52449. It appears that the only important predictor for predicting TipPercentage is Bill.

- (g) Build a neural network model with one hidden layer using the training data to predict TipPercentage. Audition number of nodes (neurons) in 1-6 and decay of 10 raised to the -2, -1.5, ..., 0.5, 1 powers. Report the estimated generalization RMSE and the RMSE on the holdout sample. Report the variable importance plot, and comment on which predictors appear most important for predicting TipPercentage.

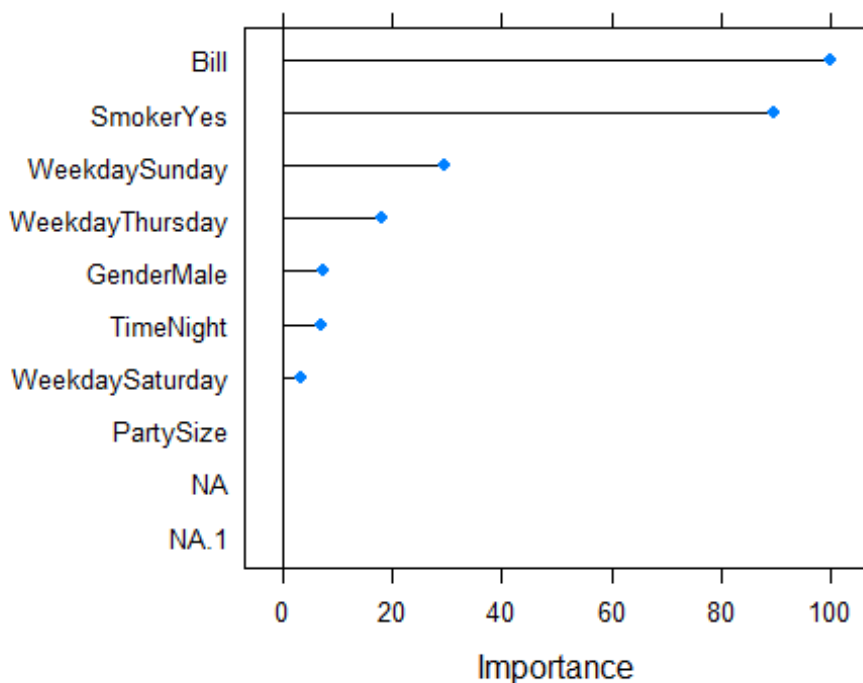
```

nnetGrid <- expand.grid(size=1:6,decay=c(10^seq(-2,1,by=0.5) ) )
set.seed(474); NNET <- train(TipPercentage~., data=TRAIN, hidden=1, method='n
net', trControl=fitControl,
                                tuneGrid=nnetGrid, preProc = c("center", "scale"
),
                                verbose=FALSE, trace=FALSE,linout=TRUE)
NNET$results[rownames(NNET$bestTune),]
##      size decay      RMSE Rsquared      MAE  RMSESD RsquaredSD      MAESD
## 35      5      10 5.835846 0.1335544 4.041775 2.142476 0.06495943 0.6147713

set.seed(474); classification.nnet <- predict(NNET,newdata=HOLDOUT)
RMSE(HOLDOUT$TipPercentage, classification.nnet)
## [1] 4.602111

IMP <- varImp(NNET)$importance
IMP$Variable <- rownames(IMP)
IMP <- IMP[order(IMP$Overall,decreasing=TRUE),]
head(IMP)
##              Overall      Variable
## Bill              100.000000      Bill
## SmokerYes          89.388789      SmokerYes
## WeekdaySunday       29.503769      WeekdaySunday
## WeekdayThursday     18.345051      WeekdayThursday
## GenderMale           7.371916      GenderMale
## TimeNight            6.960453      TimeNight
plot(varImp(NNET), top=10)

```



Response: The estimated generalization RMSE is 5.487836. The RMSE of the holdout sample is 6.078132. The predictors that appear to be most important for predicting TipPercentage are SmokerYes, Bill, and GenderMale.

(h) Is one model a compelling choice for predicting TipPercentage over the others? Why or why not? Show evidence to support your claim.

```
GLM$results
##   parameter      RMSE Rsquared      MAE  RMSESD RsquaredSD      MAESD
## 1      none 5.904779 0.1390757 4.134148 2.116873 0.07076399 0.6548951
GLMnet$results[rownames(GLMnet$bestTune),]
##   alpha  lambda      RMSE Rsquared      MAE  RMSESD RsquaredSD      MAESD
## 10      0 3.162278 5.813812 0.1323445 3.98199 2.257199 0.06256246 0.6354841
KNN$results[which(KNN$results$k == as.numeric(KNN$bestTune) ),]
##    k      RMSE Rsquared      MAE  RMSESD RsquaredSD      MAESD
## 14 14 5.742774 0.1371878 4.02328 2.165275 0.07539694 0.6360195
TREE$results[rownames(TREE$bestTune),]
##    cp      RMSE Rsquared      MAE  RMSESD RsquaredSD      MAESD
## 1 0.03589843 6.132361 0.1241258 4.132534 1.508779 0.1275679 0.3703102
FOREST$results[rownames(FOREST$bestTune),]
##   mtry      RMSE Rsquared      MAE  RMSESD RsquaredSD      MAESD
## 2      2 5.832596 0.1560506 3.979021 1.918469 0.1393561 0.5258356
GBM$results[rownames(GBM$bestTune),]
##   shrinkage interaction.depth n.minobsinnode n.trees      RMSE Rsquared
## 7          0.1                3              10      50 5.709391 0.1757444
##          MAE  RMSESD RsquaredSD      MAESD
## 7 4.050114 1.99549 0.1020458 0.5408795
NNET$results[rownames(NNET$bestTune),]
##   size decay      RMSE Rsquared      MAE  RMSESD RsquaredSD      MAESD
## 35      5      10 5.835846 0.1335544 4.041775 2.142476 0.06495943 0.6147713
```

No model is a clear winner here, as all RMSE's are well within 1 standard deviation of each other.

Problem 2. Predictive Model for Making a Purchase

The PURCHASE.csv dataset contains a small part of a customer database from a bank. Of interest is the variable Purchase, which tells us if a customer did or did not make a purchase at a major chain retailer in the following 30 days. Predictor variables include Visits (number of visits to the store in the last 90 days), Spent (how much the customer has spent in the last 90 days, PercentClose (the percentage of purchases this customer makes in general that are within 5 miles of their home address, Closest and CloseStores which details how closest the nearest store in the chain is to the customer and how many stores of that chain are within 5 miles of home).

Your task is to build a predictive model to predict Purchase (Buy/No).

Split the data into 1000 training rows with the remainder being the holdout (use set.seed(474) on the same line as the required sample command).

Use `set.seed(474)` everywhere randomness is infused in the training process.

Use 5-fold cross-validation to estimate the generalization AUC (i.e. .

```
PURCHASE <- read.csv("PURCHASE.csv")
set.seed(474); train.rows <- sample(nrow(PURCHASE),1000)
PURCHASE_TRAIN <- PURCHASE[train.rows,];
PURCHASE_HOLDOUT <- PURCHASE[-train.rows,]
fitControl <- trainControl(method="cv",number=5, classProbs=TRUE, summaryFunc
tion = twoClassSummary)
# set fitControl for estimating the generalization error. Here use AUC to tra
in and compare models.
```

- (a) Which class will the naive model classify everyone in data to? What is the estimated generalization accuracy of the this model? What is its accuracy in the holdout sample?

```
summary(PURCHASE_TRAIN$Purchase)
## Buy No
## 252 748
summary(PURCHASE_HOLDOUT$Purchase)
## Buy No
## 6640 20083
782/(218+782)
## [1] 0.782
20049/(6674+20049)
## [1] 0.7502526
```

Response: The naive model would classify everyone in the data to “No”. The estimated generalization accuracy would be 0.782. The model’s accuracy on the holdout sample would be 0.7502526.


- (b) Using the training data, train a logistic regression model to predict Purchase. Report the estimated generalization metrics of the best model, as well as its accuracy and auc on the holdout sample. Also, report the variable importance plot, and comment on which predictors appear most important for predicting Purchase.

```
set.seed(474); GLM2 <- train(Purchase~., data=PURCHASE_TRAIN, method="glm",
trControl=fitControl, preProc=c("center", "scale"
))
GLM2$results
## parameter ROC Sens Spec ROCSD SensSD Spec
SD
## 1 none 0.5827802 0.01984314 0.9866488 0.08689495 0.01414377 0.0094281
75

set.seed(474); classification.glm2 <- predict(GLM2,newdata=PURCHASE_HOLDOUT)
mean(classification.glm2==PURCHASE_HOLDOUT$Purchase)
## [1] 0.7477454
set.seed(474); roc(PURCHASE_HOLDOUT$Purchase, predict(GLM2,newdata=PURCHASE_H
OLDOUT,type="prob"))[,2])
## Setting levels: control = Buy, case = No
## Setting direction: controls < cases
```

```
##
## Call:
## roc.default(response = PURCHASE_HOLDOUT$Purchase, predictor = predict(GLM2
,      newdata = PURCHASE_HOLDOUT, type = "prob")[, 2])
##
## Data: predict(GLM2, newdata = PURCHASE_HOLDOUT, type = "prob")[, 2] in 664
0 controls (PURCHASE_HOLDOUT$Purchase Buy) < 20083 cases (PURCHASE_HOLDOUT$Pu
rchase No).
## Area under the curve: 0.625

IMP <- varImp(GLM2)$importance
IMP$Variable <- rownames(IMP)
IMP <- IMP[order(IMP$Overall, decreasing=TRUE),]
head(IMP)
##              Overall      Variable
## Visits            100.000000      Visits
## Closest           17.054036      Closest
## CloseStores       16.040762    CloseStores
## Spent              1.952926        Spent
## PercentClose       0.000000    PercentClose
plot(varImp(GLM2), top=10)
```



RESPONSE: The estimated generalization AUC is 0.5917691. The AUC for the holdout sample is 0.6214, and the accuracy of the holdout sample is 0.7457247. The predictors that appear to be most important for predicting Purchase are Visits and Spent.

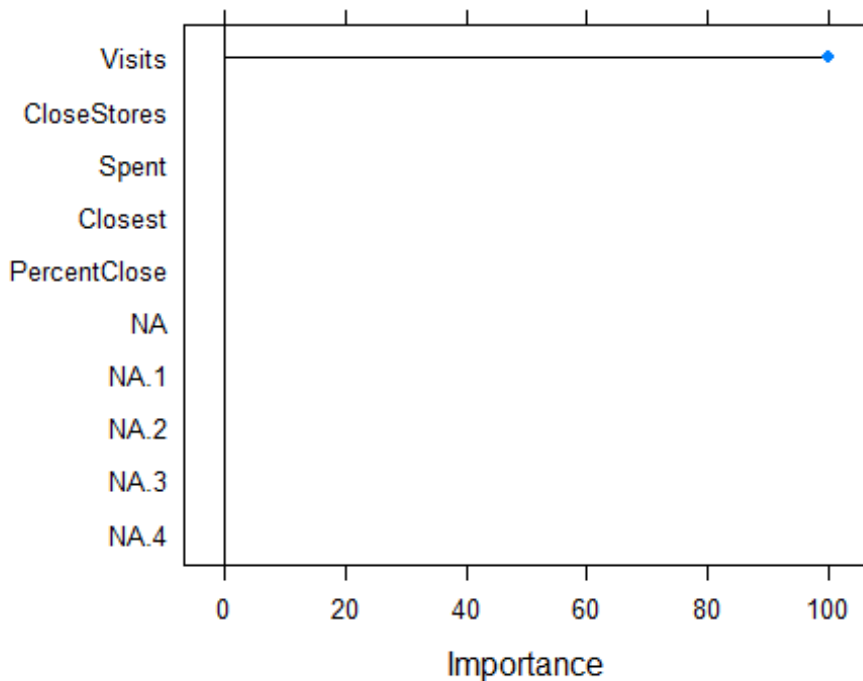
- (c) Using the training data, train a regularized logistic regression model to predict Purchase. Audition alpha values along the sequence 0, 0.1, 0.2, ..., 0.9, 1 and lambda of 10 raised to the -3, -2.5, -2, ..., 1.5, 2 powers. Report the estimated generalization metrics for the best model, as well as its accuracy and auc on the holdout samples. Also, report the variable importance plot, and comment on which predictors appear most important for predicting Purchase.

```
glmnetGrid <- expand.grid(alpha = seq(0,1, by=0.1), lambda = 10^seq(-3,2, by=
0.5))
set.seed(474); GLMnet2 <- train(Purchase~., data=PURCHASE_TRAIN, method="glmnet", tuneGrid=glmnetGrid,
                                trControl=fitControl, preProc=c("center", "scale
"))
## Warning in train.default(x, y, weights = w, ...): The metric "Accuracy" was not
## in the result set. ROC will be used instead.
GLMnet2$results[rownames(GLMnet2$bestTune),]
##   alpha      lambda      ROC Sens Spec      ROCSD      SensSD      SpecSD
## 92    0.8 0.03162278 0.6262559 0.008 0.996 0.07445817 0.01095445 0.00894427
## 2

set.seed(474); classification.glmnet2 <- predict(GLMnet2,newdata=PURCHASE_HOLDOUT)
mean(classification.glmnet2==PURCHASE_HOLDOUT$Purchase)
## [1] 0.7505894
set.seed(474); roc(PURCHASE_HOLDOUT$Purchase, predict(GLMnet2,newdata=PURCHASE_HOLDOUT,type="prob"),[,2])
## Setting levels: control = Buy, case = No
## Setting direction: controls < cases
##
## Call:
## roc.default(response = PURCHASE_HOLDOUT$Purchase, predictor = predict(GLMnet2,
## newdata = PURCHASE_HOLDOUT, type = "prob"),[, 2])
##
## Data: predict(GLMnet2, newdata = PURCHASE_HOLDOUT, type = "prob"),[, 2] in
6640 controls (PURCHASE_HOLDOUT$Purchase Buy) < 20083 cases (PURCHASE_HOLDOUT$Purchase No).
## Area under the curve: 0.6285

IMP <- varImp(GLMnet2)$importance
IMP$Variable <- rownames(IMP)
IMP <- IMP[order(IMP$Overall,decreasing=TRUE),]
head(IMP)
##           Overall      Variable
## Visits         100      Visits
## Spent           0       Spent
## PercentClose    0 PercentClose
## Closest         0       Closest
```

```
## CloseStores      0 CloseStores
plot(varImp(GLMnet2), top=10)
```



RESPONSE: The estimated generalization AUC is 0.6252065. The AUC for the holdout sample is 0.6287, and the accuracy of the holdout sample is 0.7502526. The only predictor that appears to be important for predicting Purchase is Visits.

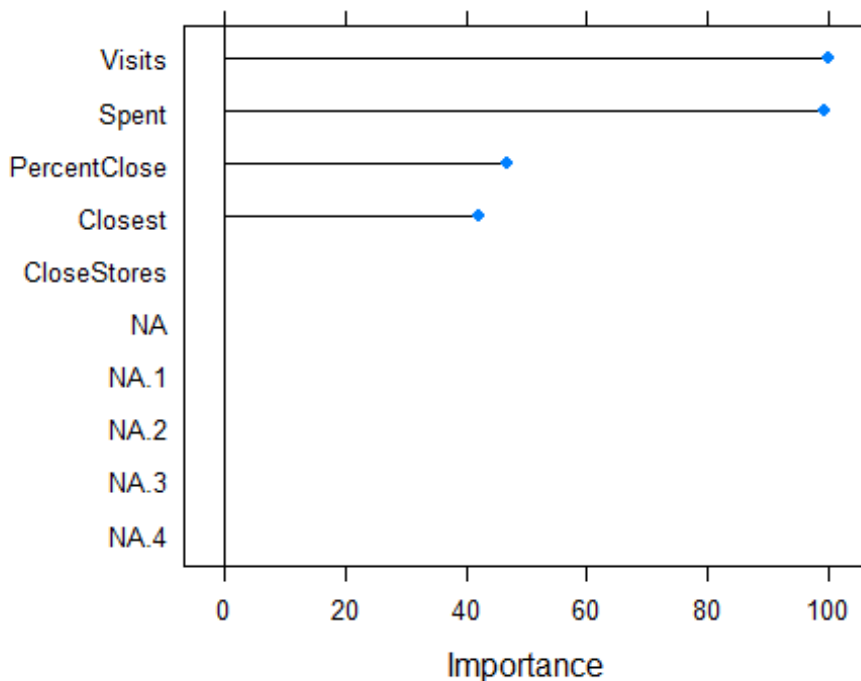
- (d) Using the training data, train a classification tree model to predict Purchase. Report the estimated generalization metrics of the best model, as well as its accuracy and auc on the holdout sample. Also, report the variable importance plot, and comment on which predictors appear most important for predicting Purchase.

```
set.seed(474); TREE2 <- train(Purchase~., data=PURCHASE_TRAIN, method="rpart",
                             rControl=fitControl,
                             preProc=c("center", "scale"))
## Warning in train.default(x, y, weights = w, ...): The metric "Accuracy" was not
## in the result set. ROC will be used instead.
TREE2$results[rownames(TREE2$bestTune),]
##          cp          ROC        Sens        Spec        ROCSD        SensSD        SpecSD
## 2 0.006613757 0.5642159 0.1705882 0.8704161 0.06042947 0.04577665 0.06030375

set.seed(474); classification.tree2 <- predict(TREE2, newdata=PURCHASE_HOLDOUT)
mean(classification.tree2==PURCHASE_HOLDOUT$Purchase)
```

```
## [1] 0.7465105
set.seed(474); roc(PURCHASE_HOLDOUT$Purchase, predict(TREE2,newdata=PURCHASE_
HOLDOUT,type="prob"))[,2])
## Setting levels: control = Buy, case = No
## Setting direction: controls < cases
##
## Call:
## roc.default(response = PURCHASE_HOLDOUT$Purchase, predictor = predict(TREE
2,
  newdata = PURCHASE_HOLDOUT, type = "prob"))[, 2])
##
## Data: predict(TREE2, newdata = PURCHASE_HOLDOUT, type = "prob"))[, 2] in 66
40 controls (PURCHASE_HOLDOUT$Purchase Buy) < 20083 cases (PURCHASE_HOLDOUT$P
urchase No).
## Area under the curve: 0.6039

IMP <- varImp(TREE2)$importance
IMP$Variable <- rownames(IMP)
IMP <- IMP[order(IMP$Overall,decreasing=TRUE),]
head(IMP)
##           Overall      Variable
## Visits      100.00000      Visits
## Spent       99.45258      Spent
## PercentClose 46.73082 PercentClose
## Closest     42.28232      Closest
## CloseStores  0.00000 CloseStores
plot(varImp(TREE2), top=10)
```



RESPONSE: The estimated generalization AUC is 0.6112599. The AUC for the holdout sample is 0.5391, and the accuracy of the holdout sample is 0.7344235. The predictors that appear to be most important for predicting Purchase are Closest, Spent, Visits, and PercentClose.


- (e) Using the training data, train a random forest model to predict Purchase. Audit two values of `mtry` of 1 and 5 (pure bagging). Report the estimated generalization metrics of the best model, as well as its accuracy and auc on the holdout sample. Also, report the variable importance plot, and comment on which predictors appear most important for predicting Purchase.

```
forestGrid <- expand.grid(mtry=c(1,5))
set.seed(474); FOREST2 <- train(Purchase~., data=PURCHASE_TRAIN, method="rf", tuneGrid=forestGrid,
                                trControl=fitControl, preProc=c("center", "scale"))
## Warning in train.default(x, y, weights = w, ...): The metric "Accuracy" was not
## in the result set. ROC will be used instead.
FOREST2$results[rownames(FOREST2$bestTune),]
##   mtry      ROC      Sens      Spec      ROCSD      SensSD      SpecSD
## 1     1 0.6364147 0.04384314 0.9679374 0.06806732 0.03849275 0.0118973

set.seed(474); classification.forest2 <- predict(FOREST2, newdata=PURCHASE_HOLDOUT)
mean(classification.forest2==PURCHASE_HOLDOUT$Purchase)
## [1] 0.7418329
set.seed(474); roc(PURCHASE_HOLDOUT$Purchase, predict(FOREST2, newdata=PURCHASE_HOLDOUT, type="prob"), 2)
## Setting levels: control = Buy, case = No
## Setting direction: controls < cases
##
## Call:
## roc.default(response = PURCHASE_HOLDOUT$Purchase, predictor = predict(FOREST2,
##                               newdata = PURCHASE_HOLDOUT, type = "prob"), 2)
##
## Data: predict(FOREST2, newdata = PURCHASE_HOLDOUT, type = "prob"), 2 in
6640 controls (PURCHASE_HOLDOUT$Purchase Buy) < 20083 cases (PURCHASE_HOLDOUT$Purchase No).
## Area under the curve: 0.608

IMP <- varImp(FOREST2)$importance
IMP$Variable <- rownames(IMP)
IMP <- IMP[order(IMP$Overall, decreasing=TRUE),]
head(IMP)
##           Overall      Variable
## Spent      100.00000      Spent
## PercentClose 95.90545 PercentClose
## Closest     78.34374      Closest
```

```
## Visits      36.87581      Visits
## CloseStores 0.00000 CloseStores
plot(varImp(FOREST2), top=10)
```



RESPONSE: The estimated generalization AUC is 0.5619115. The AUC for the holdout sample is 0.6062, and the accuracy of the holdout sample is 0.7476705. The predictors that appear to be most important for predicting Purchase are Closest, Spent, Visits, and PercentClose.

- (f) Using the training data, train a gradient boosted tree model to predict Purchase. Report the estimated generalization metrics of the best model, as well as its accuracy and auc on the holdout sample. Also, report the variable importance plot, and comment on which predictors appear most important for predicting Purchase.

```
set.seed(474); GBM2 <- train(Purchase~., data=PURCHASE_TRAIN, method="gbm",
                             trControl=fitControl, preProc=c("center", "scale"),
                             verbose=FALSE)
## Warning in train.default(x, y, weights = w, ...): The metric "Accuracy" was not
## in the result set. ROC will be used instead.
GBM2$results[rownames(GBM2$bestTune),]
## shrinkage interaction.depth n.minobsinnode n.trees      ROC  Sens  Spec
## 1      0.1              1              10      50 0.6320067 0.012 0.996
##      ROCSD      SensSD      SpecSD
## 1 0.07886722 0.01788854 0.008944272

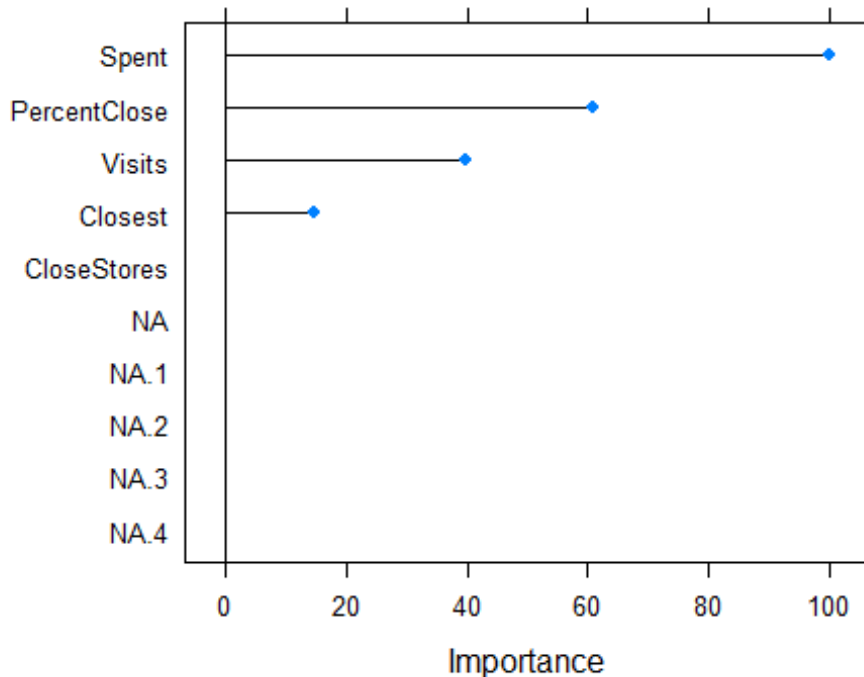
set.seed(474); classification.gbm2 <- predict(GBM2, newdata=PURCHASE_HOLDOUT)
```

```

mean(classification.gbm2==PURCHASE_HOLDOUT$Purchase)
## [1] 0.7509636
set.seed(474); roc(PURCHASE_HOLDOUT$Purchase, predict(GBM2,newdata=PURCHASE_H
OLDOUT,type="prob"),2))
## Setting levels: control = Buy, case = No
## Setting direction: controls < cases
##
## Call:
## roc.default(response = PURCHASE_HOLDOUT$Purchase, predictor = predict(GBM2
, newdata = PURCHASE_HOLDOUT, type = "prob"), 2))
##
## Data: predict(GBM2, newdata = PURCHASE_HOLDOUT, type = "prob"), 2] in 664
0 controls (PURCHASE_HOLDOUT$Purchase Buy) < 20083 cases (PURCHASE_HOLDOUT$Pu
rchase No).
## Area under the curve: 0.6168

IMP <-varImp(GBM2)$importance
IMP$Variable <-rownames(IMP)
IMP <- IMP[order(IMP$Overall,decreasing=TRUE),]
head(IMP)
##           Overall      Variable
## Spent          100.00000      Spent
## PercentClose   60.81310 PercentClose
## Visits         39.77526      Visits
## Closest        14.67604      Closest
## CloseStores    0.00000 CloseStores
plot(varImp(GBM2), top=10)

```



RESPONSE: The estimated generalization AUC is 0.5780394. The AUC for the holdout sample is 0.6087, and the accuracy of the holdout sample is 0.749841. The predictors that appear to be most important for predicting Purchase are Spent, Visits, PercentClose, and Closest.

- (g) Using the training data, train a support vector machine with a radial basis kernel to predict Purchase. Report the estimated generalization metrics of the best model, as well as its accuracy and auc on the holdout samples. Also, report the variable importance plot, and comment on which predictors appear most important for predicting Purchase.

```
set.seed(474); SVMradial <- train(Purchase~., data=PURCHASE_TRAIN, method="svmRadial",
                                trControl=fitControl, verbose=FALSE, preProc=c("center", "scale"))
## Warning in train.default(x, y, weights = w, ...): The metric "Accuracy" was not
## in the result set. ROC will be used instead.
SVMradial$results[rownames(SVMradial$bestTune),]
##      sigma C      ROC Sens      Spec      ROCSD      SensSD      SpecSD
## 3 0.4511284 1 0.5718481 0.004 0.9986667 0.04128938 0.008944272 0.002981424


set.seed(474); classification.svm <- predict(SVMradial, newdata=PURCHASE_HOLDOUT)
mean(classification.svm==PURCHASE_HOLDOUT$Purchase)
## [1] 0.7511133
set.seed(474); roc(PURCHASE_HOLDOUT$Purchase, predict(SVMradial, newdata=PURCH
```

```

ASE_HOLDOUT,type="prob")[,2])
## Setting levels: control = Buy, case = No
## Setting direction: controls < cases
##
## Call:
## roc.default(response = PURCHASE_HOLDOUT$Purchase, predictor = predict(SVMr
adial,      newdata = PURCHASE_HOLDOUT, type = "prob")[, 2])
##
## Data: predict(SVMradial, newdata = PURCHASE_HOLDOUT, type = "prob")[, 2] i
n 6640 controls (PURCHASE_HOLDOUT$Purchase Buy) < 20083 cases (PURCHASE_HOLD
OUT$Purchase No).
## Area under the curve: 0.5445

plot(varImp(SVMradial), top=10)

```



RESPONSE: The estimated generalization AUC is 0.5066999. The AUC for the holdout sample is 0.52, and the accuracy of the holdout sample is 0.7502526. The predictors that appear to be most important for predicting Purchase are Visits and Spent.

- (h) Using the training data, train a neural network model with one hidden layer to predict Purchase. Audition number of nodes in 1-6 and decay of 10 raised to the -2, -1.5, ..., 0.5, 1 powers. Report the estimated generalization metrics of the best model, as well as the accuracy and auc on the holdout sample. Also, report the variable importance plot, and comment on which predictors appear most important for predicting Purchase.

```

nnetGrid <- expand.grid(size=1:6,decay=c(10^seq(-2,1,by=0.5)) )
set.seed(474); NNET2 <- train(Purchase~., data=PURCHASE_TRAIN, hidden=1, meth

```

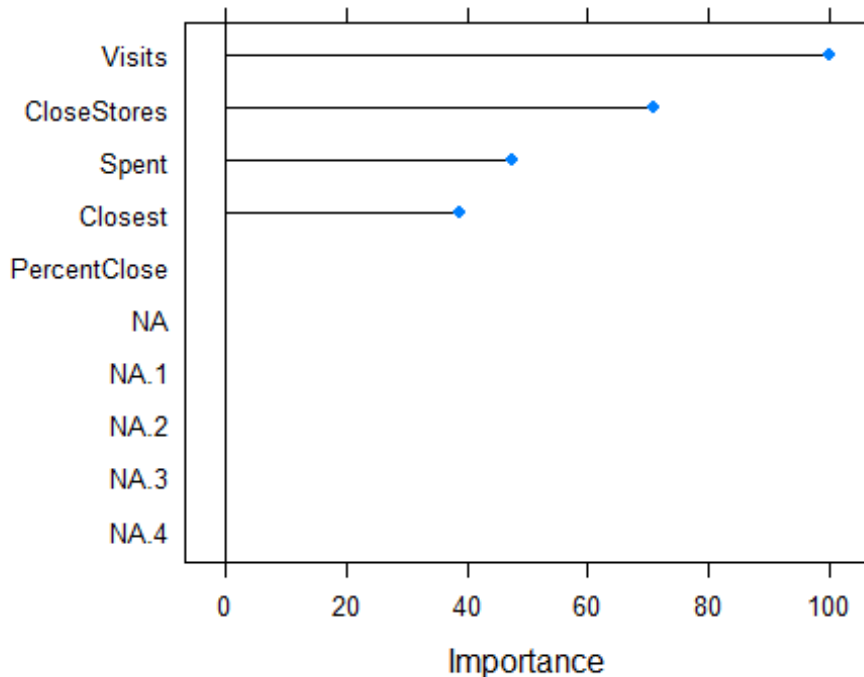
```

od='nnet', trControl=fitControl,
                                tuneGrid=nnetGrid, preProc = c("center", "scale"
),
                                verbose=FALSE, trace=FALSE, linout=FALSE)
## Warning in train.default(x, y, weights = w, ...): The metric "Accuracy" was not
## in the result set. ROC will be used instead.
NNET2$results[rownames(NNET2$bestTune),]
##      size      decay      ROC      Sens      Spec      ROCSD      SensSD
## 16      3 0.03162278 0.6182185 0.02792157 0.9799821 0.04474965 0.02283855
##      SpecSD
## 16 0.016988

set.seed(474); classification.nnet2 <- predict(NNET2, newdata=PURCHASE_HOLDOUT
)
mean(classification.nnet2==PURCHASE_HOLDOUT$Purchase)
## [1] 0.7456498
set.seed(474); roc(PURCHASE_HOLDOUT$Purchase, predict(NNET2, newdata=PURCHASE_
HOLDOUT, type="prob"),[,2])
## Setting levels: control = Buy, case = No
## Setting direction: controls < cases
##
## Call:
## roc.default(response = PURCHASE_HOLDOUT$Purchase, predictor = predict(NNET
2,      newdata = PURCHASE_HOLDOUT, type = "prob"),[, 2])
##
## Data: predict(NNET2, newdata = PURCHASE_HOLDOUT, type = "prob"),[, 2] in 66
40 controls (PURCHASE_HOLDOUT$Purchase Buy) < 20083 cases (PURCHASE_HOLDOUT$P
urchase No).
## Area under the curve: 0.6098

IMP <- varImp(NNET2)$importance
IMP$Variable <- rownames(IMP)
IMP <- IMP[order(IMP$Overall, decreasing=TRUE),]
head(IMP)
##           Overall      Variable
## Visits      100.00000      Visits
## CloseStores  71.03920 CloseStores
## Spent        47.53500      Spent
## Closest      38.73402      Closest
## PercentClose  0.00000 PercentClose
plot(varImp(NNET2), top=10)

```



RESPONSE: The estimated generalization AUC is 0.6160569. The AUC for the holdout sample is 0.6233, and the accuracy of the holdout sample is 0.7502526. The only predictor that appears to be important for predicting Purchase is Visits.

- (i) Is one model a compelling choice for predicting Purchase over the others? Why or why not? Show evidence to support your claim.

```
GLM2$results
##   parameter      ROC      Sens      Spec      ROCSD      SensSD      SpecSD
## 1      none 0.5827802 0.01984314 0.9866488 0.08689495 0.01414377 0.0094281
75
GLMnet2$results[rownames(GLMnet2$bestTune),]
##   alpha      lambda      ROC      Sens      Spec      ROCSD      SensSD      SpecSD
## 92   0.8 0.03162278 0.6262559 0.008 0.996 0.07445817 0.01095445 0.00894427
2
TREE2$results[rownames(TREE2$bestTune),]
##   cp      ROC      Sens      Spec      ROCSD      SensSD      SpecSD
## 2 0.006613757 0.5642159 0.1705882 0.8704161 0.06042947 0.04577665 0.060303
75
FOREST2$results[rownames(FOREST2$bestTune),]
##   mtry      ROC      Sens      Spec      ROCSD      SensSD      SpecSD
## 1    1 0.6364147 0.04384314 0.9679374 0.06806732 0.03849275 0.0118973
GBM2$results[rownames(GBM2$bestTune),]
##   shrinkage interaction.depth n.minobsinnode n.trees      ROC      Sens      Spec
```

```
## 1      0.1      1      10      50 0.6320067 0.012 0.996
##      ROCSD      SensSD      SpecSD
## 1 0.07886722 0.01788854 0.008944272
SVMradial$results[rownames(SVMradial$bestTune),]
##      sigma C      ROC      Sens      Spec      ROCSD      SensSD      SpecSD
## 3 0.4511284 1 0.5718481 0.004 0.9986667 0.04128938 0.008944272 0.002981424
NNET2$results[rownames(NNET2$bestTune),]
##      size      decay      ROC      Sens      Spec      ROCSD      SensSD
## 16      3 0.03162278 0.6182185 0.02792157 0.9799821 0.04474965 0.02283855
##      SpecSD
## 16 0.016988
```

RESPONSE: Again, there is no clear winner here as most of the models' ROC's are within 1 standard deviation of each other. However, we can rule out the SVMradial model as its ROC is well below of 1 standard deviation of most models. Also, it should be noted that it is never a good idea to use AUC to evaluate SVM models. We can also likely rule out the random forest model as it is below of 1 standard deviation of a couple of models. That being said, I do not see a definitive "best model" here.